

Selective Coloring

Francisco Noya (fnoya2@illinois.edu)
Mesay Taye (mesayst2@illinois.edu)

2022-05-01

In this project we implemented a tool to selectively and automatically colorize the foreground of black and white photographs. The tool consists of two main algorithms. First, a GAN network colorizes the entire photograph using transfer learning. Second, a segmentation network segments the foreground from the background. Lastly, it blends the colorized foreground with the original background to produce a distinctive image.

In this notebook the Colorization network is trained and validated.

Colorization Network

Approach

For the colorization task we worked on the **Lab** colorspace. In this space, the **L** channel contains the luminance or intensity information, while the **ab** channels contain the color information. Therefore, a neural network can be trained with the **L** channel of regular color images as input. Its predictions will be “fake” **ab** channels and its loss will be calculated with the “real” **ab** channels.

After a short bibliographic review we found that although traditional convolutional neural networks (CNNs) could produce results almost indistinguishable from real color photos (Zhang et al., 2016), *Generative Adversarial Networks* or GANs (Goodfellow et al., 2014) were the most proper approach for this kind of problem. This network architecture contains two modules, a Generator and a Discriminator. Both models are trained in parallel. The objective of the generator is to produce outputs similar enough to the ground truth that can fool the discriminator. The discriminator’s objective is to properly tell the ground truth from the discriminator output.

Since training this kind of networks requires large datasets and computing time, we decided to use pretrained models that have been used for other tasks such as object classification. We followed a tutorial inspired by the *pix2pix* paper (Isola et al., 2016) but instead of training a naïve *UNet* as the generator, it used a *ResNet18* network as the generator (*Colorizing Black & White Images with u-Net and Conditional GAN — a Tutorial*, n.d.). Similar to *pix2pix* we used a

patch discriminator that splits the image in 26 square patches (depending on the image size) and produces a *real* or *fake* prediction for each of them.

In order to try different approaches, we decided to use *Transformers* in place of the discriminator and the generator. *Transformers* are a special architectures of DNN that make extensive use of attention mechanisms (Vaswani et al., 2017). Because of their ability to have larger receptive fields compared to convolutional neural networks (CNNs) that allow tracking long-range dependencies within an image, these attention based architectures have proven very effective in image processing tasks and gave rise to Visual Transformers or **ViT** (Dosovitskiy et al., 2020). We tried different architectures with ViTs generators or discriminators and measured a range of metrics for each of them.

```
!pip install fastai

import os
import glob
import time
import numpy as np
from PIL import Image
from pathlib import Path
from tqdm.notebook import tqdm
import matplotlib.pyplot as plt
from skimage.color import rgb2lab, lab2rgb

import torch
import cv2
from torch import nn, optim
from torchvision import transforms
from torchvision.utils import make_grid
from torch.utils.data import Dataset, DataLoader

from fastai.vision.learner import create_body
from torchvision.models.resnet import resnet18
from torchvision.models import vit_b_16 as visiontransformer
from fastai.vision.models.unet import DynamicUnet
from fastai.data.external import untar_data, URLs

device = torch.device("cuda" if torch.cuda.is_available() else
                     "cpu")
print (device)

cpu
```

Dataset

For training and validation purposes we used a subset of the COCO dataset of images (Lin et al., 2014) that is provided by the FastAI framework (*Fast.ai*, n.d.). We downloaded 10.000 images from this dataset and randomly splitted them into two sets: a training set with 8.000 images and a validation set with 2.000 images.

```
coco_path = untar_data(URLs.COCO_SAMPLE)
coco_path = str(coco_path) + "/train_sample"

path = coco_path

paths = glob.glob(path + "/*.jpg")
np.random.seed(123)
paths_subset = np.random.choice(paths, 10_000, replace=False) #
    ↪ choosing 10000 images randomly
rand_idxs = np.random.permutation(10_000)
train_idxs = rand_idxs[:8000] # choosing the first 8000 as
    ↪ training set
val_idxs = rand_idxs[8000:] # choosing last 2000 as validation
    ↪ set
train_paths = paths_subset[train_idxs]
val_paths = paths_subset[val_idxs]
print(len(train_paths), len(val_paths))

fig, axes = plt.subplots(4, 4, figsize=(10, 10))
fig.suptitle('Sample images from COCO dataset')
for ax, img_path in zip(axes.flatten(), train_paths):
    ax.imshow(Image.open(img_path))
    ax.axis("off")
```

8000 2000

Then we resized the images so that they have manageable dimensions that allow feeding into the different network architectures without requiring extremely high computational resources or long times. Similar to (Isola et al., 2016) data augmentation was achieved by flipping the images horizontally (this is only done for the training set). We used 16 images on each batch that goes through the network. Each image was converted to the **Lab** colorspace and the channels adjusted float values between -1 and 1.

```
SIZE = 224
class ColorizationDataset(Dataset):
    def __init__(self, paths, split='train'):
        if split == 'train':
            self.transforms = transforms.Compose([
```

Sample images from COCO dataset

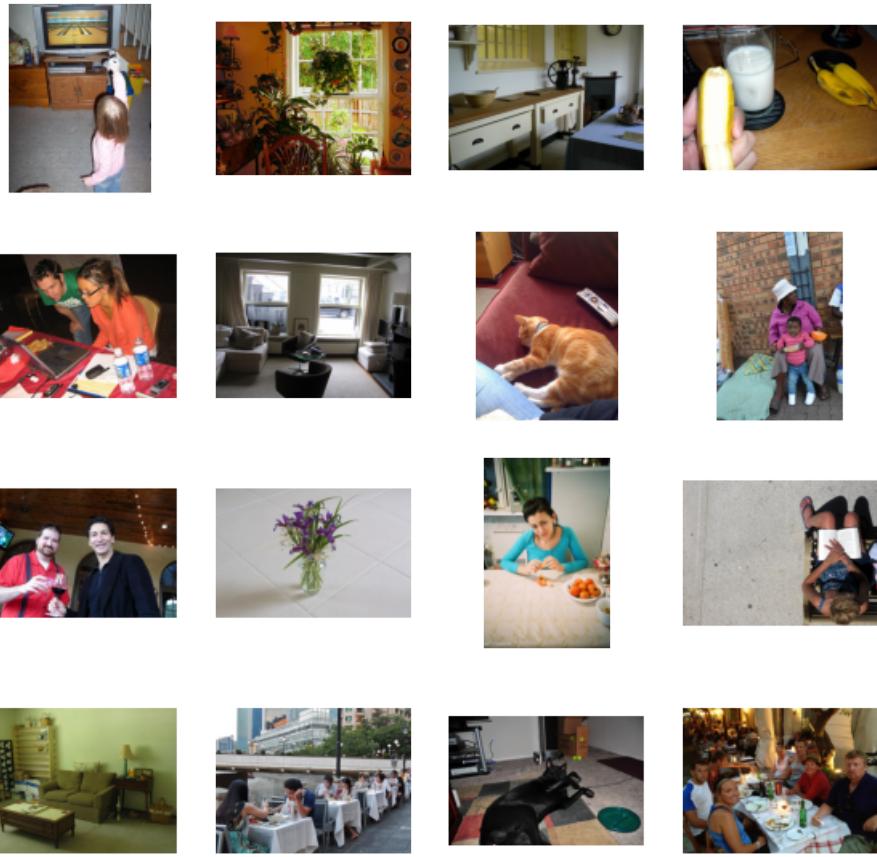


Figure 1: png

```

        transforms.Resize((SIZE, SIZE),
↪  Image.Resampling.BICUBIC),
        transforms.RandomHorizontalFlip(), # A little
↪  data augmentation!
    ])
elif split == 'val':
    self.transforms = transforms.Resize((SIZE, SIZE),
↪  Image.Resampling.BICUBIC)

self.split = split
self.size = SIZE
self.paths = paths

def __getitem__(self, idx):
    img = Image.open(self.paths[idx]).convert("RGB")
    img = self.transforms(img)
    img = np.array(img)
    img_lab = rgb2lab(img).astype("float32") # Converting RGB
↪  to L*a*b
    img_lab = transforms.ToTensor()(img_lab)
    L = img_lab[[0], ...] / 50. - 1. # Between -1 and 1
    ab = img_lab[[1, 2], ...] / 110. # Between -1 and 1

    return {'L': L, 'ab': ab}

def __len__(self):
    return len(self.paths)

def make_dataloaders(batch_size=16, n_workers=2, pin_memory=True,
↪  **kwargs): # A handy function to make our dataloaders
    if (os.name=="nt"):
        n_workers = 0 # In Windows, n_workers should be 0
    dataset = ColorizationDataset(**kwargs)
    dataloader = DataLoader(dataset, batch_size=batch_size,
↪  num_workers=n_workers,
                           pin_memory=pin_memory)
    return dataloader

train_dl = make_dataloaders(paths=train_paths, split='train')
val_dl = make_dataloaders(paths=val_paths, split='val')

data = next(iter(train_dl))
Ls, abs_ = data['L'], data['ab']
print(Ls.shape, abs_.shape)

```

```
torch.Size([16, 1, 224, 224]) torch.Size([16, 2, 224, 224])
```

Loss functions

The loss function of the discriminator for each image is the binary cross entropy between the predictions and the ground truth: *real* if the real **ab** channels were fed into the discriminator or *fake* if the generated **ab** channels were used instead. The loss function of the generator was the combination of the L1 error and the loss function of the discriminator as it were *real ab* channels. The intuition behind is that the generator “wins” every time it fools the discriminator into assigning *real* predictions to its outputs.

```
class GANLoss(nn.Module):
    def __init__(self, gan_mode='vanilla', real_label=1.0,
                 fake_label=0.0):
        super().__init__()
        self.register_buffer('real_label',
                            torch.tensor(real_label))
        self.register_buffer('fake_label',
                            torch.tensor(fake_label))
        if gan_mode == 'vanilla':
            self.loss = nn.BCEWithLogitsLoss()
        elif gan_mode == 'lsgan':
            self.loss = nn.MSELoss()

    def get_labels(self, preds, target_is_real):
        if target_is_real:
            labels = self.real_label
        else:
            labels = self.fake_label
        return labels.expand_as(preds)

    def __call__(self, preds, target_is_real):
        labels = self.get_labels(preds, target_is_real)
        loss = self.loss(preds, labels)
        return loss

    def init_weights(net, init='norm', gain=0.02):

        def init_func(m):
            classname = m.__class__.__name__
            if hasattr(m, 'weight') and 'Conv' in classname:
                if init == 'norm':
                    nn.init.normal_(m.weight.data, mean=0.0,
                                   std=gain)
                elif init == 'xavier':
```

```

        nn.init.xavier_normal_(m.weight.data, gain=gain)
    elif init == 'kaiming':
        nn.init.kaiming_normal_(m.weight.data, a=0,
→      mode='fan_in')

        if hasattr(m, 'bias') and m.bias is not None:
            nn.init.constant_(m.bias.data, 0.0)
    elif 'BatchNorm2d' in classname:
        nn.init.normal_(m.weight.data, 1., gain)
        nn.init.constant_(m.bias.data, 0.)

    net.apply(init_func)
    print(f"model initialized with {init} initialization")
    return net

def init_model(model, device):
    model = model.to(device)
    model = init_weights(model)
    return model

class AverageMeter:
    def __init__(self):
        self.reset()

    def reset(self):
        self.count, self.avg, self.sum = [0.] * 3

    def update(self, val, count=1):
        self.count += count
        self.sum += count * val
        self.avg = self.sum / self.count

    def create_loss_meters():
        loss_D_fake = AverageMeter()
        loss_D_real = AverageMeter()
        loss_D = AverageMeter()
        loss_G_GAN = AverageMeter()
        loss_G_L1 = AverageMeter()
        loss_G = AverageMeter()

        return {'loss_D_fake': loss_D_fake,
                'loss_D_real': loss_D_real,
                'loss_D': loss_D,
                'loss_G_GAN': loss_G_GAN,
                'loss_G_L1': loss_G_L1,

```

```

        'loss_G': loss_G}

def update_losses(model, loss_meter_dict, count):
    for loss_name, loss_meter in loss_meter_dict.items():
        loss = getattr(model, loss_name)
        loss_meter.update(loss.item(), count=count)

def lab_to_rgb(L, ab):
    """
    Takes a batch of images
    """

    L = (L + 1.) * 50.
    ab = ab * 110.
    Lab = torch.cat([L, ab], dim=1).permute(0, 2, 3,
→ 1).cpu().numpy()
    rgb_imgs = []
    for img in Lab:
        img_rgb = lab2rgb(img)
        rgb_imgs.append(img_rgb)
    return np.stack(rgb_imgs, axis=0)

def model_eval(model, data):
    model.net_G.eval()
    with torch.no_grad():
        model.setup_input(data)
        model.forward()
    model.net_G.train()
    fake_color = model.fake_color.detach()
    real_color = model.ab
    L = model.L
    fake_imgs = lab_to_rgb(L, fake_color)
    real_imgs = lab_to_rgb(L, real_color)
    return fake_imgs, real_imgs

def visualize(model, data, save=True):
    model.net_G.eval()
    with torch.no_grad():
        model.setup_input(data)
        model.forward()
    model.net_G.train()
    fake_color = model.fake_color.detach()
    real_color = model.ab
    L = model.L
    fake_imgs = lab_to_rgb(L, fake_color)

```

```

real_imgs = lab_to_rgb(L, real_color)
fig = plt.figure(figsize=(15, 8))
for i in range(5):
    ax = plt.subplot(3, 5, i + 1)
    ax.imshow(L[i][0].cpu(), cmap='gray')
    ax.axis("off")
    ax = plt.subplot(3, 5, i + 1 + 5)
    ax.imshow(fake_imgs[i])
    ax.axis("off")
    ax = plt.subplot(3, 5, i + 1 + 10)
    ax.imshow(real_imgs[i])
    ax.axis("off")
plt.show()
if save:
    fig.savefig(f"output/colorization_{time.time()}.png")
return fake_imgs, real_imgs

def log_results(loss_meter_dict):
    for loss_name, loss_meter in loss_meter_dict.items():
        print(f"{loss_name}: {loss_meter.avg:.5f}")

```

Model

```

class MainModel(nn.Module):
    def __init__(self, net_G=None, net_D=None, use_ViT_gen =
    ↪ False, lr_G=2e-4, lr_D=2e-4,
                beta1=0.5, beta2=0.999, lambda_L1=100.):
        super().__init__()

        self.device = torch.device("cuda" if
        ↪ torch.cuda.is_available() else "cpu")
        self.lambda_L1 = lambda_L1
        self.use_ViT_gen = use_ViT_gen

        if net_G is None:
            raise NotImplementedError
        else:
            self.net_G = net_G.to(self.device)

        if net_D is None:
            self.net_D = init_model(PatchDiscriminator(input_c=3,
            ↪ n_down=3, num_filters=64), self.device)
        else:
            self.net_D = net_D.to(self.device)

```

```

#self.GANcriterion =
    ↵ GANLoss(gan_mode='vanilla').to(self.device) #
    ↵ Original BCE Loss
self.GANcriterion =
    ↵ GANLoss(gan_mode='lsgan').to(self.device) # Final
    ↵ improvement with Least Square Error loss
self.L1criterion = nn.L1Loss()
self.opt_G = optim.Adam(self.net_G.parameters(), lr=lr_G,
    ↵ betas=(beta1, beta2))
self.opt_D = optim.Adam(self.net_D.parameters(), lr=lr_D,
    ↵ betas=(beta1, beta2))

def set_requires_grad(self, model, requires_grad=True):
    for p in model.parameters():
        p.requires_grad = requires_grad

def setup_input(self, data):
    self.L = data['L'].to(self.device)
    self.ab = data['ab'].to(self.device)

def forward(self):
    if (self.use_ViT_gen == True):
        outputs = self.net_G(self.L.repeat(1,3,1,1))
        self.fake_color = outputs.logits
    else:
        self.fake_color = self.net_G(self.L)

def backward_D(self):
    fake_image = torch.cat([self.L, self.fake_color], dim=1)
    fake_preds = self.net_D(fake_image.detach())
    self.loss_D_fake = self.GANcriterion(fake_preds, False)
    real_image = torch.cat([self.L, self.ab], dim=1)
    real_preds = self.net_D(real_image)
    self.loss_D_real = self.GANcriterion(real_preds, True)
    self.loss_D = (self.loss_D_fake + self.loss_D_real) * 0.5
    self.loss_D.backward()

def backward_G(self):
    fake_image = torch.cat([self.L, self.fake_color], dim=1)
    fake_preds = self.net_D(fake_image)
    self.loss_G_GAN = self.GANcriterion(fake_preds, True)
    self.loss_G_L1 = self.L1criterion(self.fake_color,
        ↵ self.ab) * self.lambda_L1
    self.loss_G = self.loss_G_GAN + self.loss_G_L1

```

```

        self.loss_G.backward()

    def optimize(self):
        self.forward()
        self.net_D.train()
        self.set_requires_grad(self.net_D, True)
        self.opt_D.zero_grad()
        self.backward_D()
        self.opt_D.step()

        self.net_G.train()
        self.set_requires_grad(self.net_D, False)
        self.opt_G.zero_grad()
        self.backward_G()
        self.opt_G.step()

class PatchDiscriminator(nn.Module):
    def __init__(self, input_c, num_filters=64, n_down=3):
        super().__init__()
        model = [self.get_layers(input_c, num_filters,
→ norm=False)]
        model += [self.get_layers(num_filters * 2 ** i,
→ num_filters * 2 ** (i + 1), s=1 if i == (n_down-1) else 2
                    for i in range(n_down)) # the 'if'
→ statement is taking care of not
→ using
                    # stride of 2
→ for the
→ last block
→ in this
→ loop
        model += [self.get_layers(num_filters * 2 ** n_down, 1,
→ s=1, norm=False, act=False)] # Make sure to not use
→ normalization or

        self.model = nn.Sequential(*model)
→

```

```

def get_layers(self, ni, nf, k=4, s=2, p=1, norm=True,
← act=True): # when needing to make some repetitive blocks
← of layers,
    layers = [nn.Conv2d(ni, nf, k, s, p, bias=not norm)]
← # it's always helpful to make a separate method for that
← purpose
    if norm: layers += [nn.BatchNorm2d(nf)]
    if act: layers += [nn.LeakyReLU(0.2, True)]
    return nn.Sequential(*layers)

def forward(self, x):
    return self.model(x)

def build_res_unet(n_input=1, n_output=2, size=256):
    device = torch.device("cuda" if torch.cuda.is_available()
← else "cpu")
    body = create_body(resnet18, pretrained=True, n_in=n_input,
← cut=-2)
    net_G = DynamicUnet(body, n_output, (size, size)).to(device)
    return net_G

def build_visiontransformer(n_output=900, size=256):
    net_d = visiontransformer(image_size=size)
    net_d.heads = nn.Linear(768, n_output)
    torch.nn.init.xavier_uniform_(net_d.heads.weight)
    return net_d

# Build transformer based generator
# https://huggingface.co/docs/transformers/model_doc/vit

from transformers import ViTForMaskedImageModeling, ViTConfig

def build_VTi_generator():
    device = torch.device("cuda" if torch.cuda.is_available()
← else "cpu")
    model =
    ViTForMaskedImageModeling.from_pretrained("google/vit-base-patch16-224-in21k")
    # config = ViTConfig(num_channels=1)
    # model = ViTForMaskedImageModeling(config)
    # model.decoder = nn.Sequential(nn.Conv2d(768, 512,
    ← kernel_size=(1, 1), stride=(1, 1)),
    ← nn.PixelShuffle(upscale_factor=16)) ## version 1 channel
    model.decoder = nn.Sequential(nn.Conv2d(768, 768,
    ← kernel_size=3, stride=1, padding=1),

```

```

                nn.ReLU(inplace=True),
                nn.Conv2d(768, 768,
↪  kernel_size=3, stride=1, padding=1), \
                nn.ReLU(inplace=True),
                nn.Conv2d(768, 512,
↪  kernel_size=3, stride=1, padding=1), \
↪  nn.PixelShuffle(upscale_factor=16))
model = model.to(device)
return model

def pretrain_generator(net_G, train_dl, opt, criterion, epochs):
    for e in range(epochs):
        loss_meter = AverageMeter()
        for data in tqdm(train_dl):
            L, ab = data['L'].to(device), data['ab'].to(device)
            preds = net_G(L.repeat(1,3,1,1))
            preds = preds.logits
            loss = criterion(preds, ab)
            opt.zero_grad()
            loss.backward()
            opt.step()

            loss_meter.update(loss.item(), L.size(0))

            print(f"Epoch {e + 1}/{epochs}")
            print(f'L1 Loss: {loss_meter.avg:.5f}')
            #torch.save(net_G.state_dict(),
↪  'models/net_G_resnet18_model-' + str(e) +'.pt')

def train_model(model, train_dl, epochs, display_every=200,
↪  first_epoch=0):
    data_val = next(iter(val_dl)) # getting a batch for
    ↪  visualizing the model output after fixed intervals
    for e in range(first_epoch, epochs):
        loss_meter_dict = create_loss_meters() # function
        ↪  returning a dictionary of objects to
        i = 0 # log the losses
        ↪  of the complete network
        for data in tqdm(train_dl):
            model.setup_input(data)
            model.optimize()
            update_losses(model, loss_meter_dict,
↪  count=data['L'].size(0)) # function updating the log objects
            i += 1

```

```

        if i % display_every == 0:
            print(f"\nEpoch {e+1}/{epochs}")
            print(f"Iteration {i}/{len(train_dl)}")
            log_results(loss_meter_dict) # function to print
    ↵  out the losses
            visualize(model, data_val, save=False) # function
    ↵  displaying the model's outputs
            torch.save(model.state_dict(),
    ↵  'models/model6-3channel-ViT.pt')

```

Metrics

To assess the different networks architectures we selected a set of metrics for regression models:

- Correlation coefficient R squared
- Explained variance
- Mean absolute error
- Median absolute error
- Mean squared error

We calculated all these metrics for each one of the **ab** channels.

```

from sklearn.metrics import r2_score, explained_variance_score,
    ↵  mean_absolute_error, mean_squared_error,
    ↵  median_absolute_error
import pandas as pd

def get_metrics(model, dl):
    test_iter = iter(dl)
    [fa,fb,ra,rb] = [np.zeros((1,SIZE**2))] * 4

    for data in tqdm(test_iter):
        fake_imgs, real_imgs = model_eval(model, data)
        fab =
    ↵  np.reshape(fake_imgs[:,...,1],(fake_imgs.shape[0],-1))
        rab =
    ↵  np.reshape(real_imgs[:,...,1],(real_imgs.shape[0],-1))
        fbb =
    ↵  np.reshape(fake_imgs[:,...,2],(fake_imgs.shape[0],-1))
        rbb =
    ↵  np.reshape(real_imgs[:,...,2],(real_imgs.shape[0],-1))
        fa = np.concatenate((fa,fab))
        fb = np.concatenate((fb,fbb))
        ra = np.concatenate((ra,rab))

```

```

rb = np.concatenate((rb,rbb))

table=[]
table.append(["R-square", r2_score(fa, ra), r2_score(fb,
↪ rb)])
table.append(["Explained variance",
↪ explained_variance_score(fa, ra),
↪ explained_variance_score(fb, rb)])
table.append(["Mean absolute error", mean_absolute_error(fa,
↪ ra), mean_absolute_error(fb, rb)])
table.append(["Median absolute error",
↪ median_absolute_error(fa, ra), median_absolute_error(fb,
↪ rb)])
table.append(["Mean squared error", mean_squared_error(fa,
↪ ra), mean_squared_error(fb, rb)])
table.append(["Sample size", fa.shape[0], fb.shape[0]])

df = pd.DataFrame(np.array(table),columns =
↪ ['Metric','a-channel','b-channel'])
return df

```

First approach: *ResNet18* generator

One of the challenges of GANs is that, at the beginning of the training, the task of the discriminator is much easier than that of the generator because the generated outputs are very different from the real ones. In this situation, the discriminator learns so much faster and gives no time to the generator to adapt. To avoid this, we gave the generator a *head start* by training it alone (without the generator) for 20 epochs with a L1 loss function and saving its weights.

```

# Pretraining of ResNet18 Generator

net_G = build_res_unet(n_input=1, n_output=2, size=256)
opt = optim.Adam(net_G.parameters(), lr=1e-4)
criterion = nn.L1Loss()
pretrain_generator(net_G, train_dl, opt, criterion, 20)
torch.save(net_G.state_dict(),
↪ "models/net_G_resnet18_model-19.pt")

```

After that we started the parallel training of the generator and the patch discriminator for another 20 epochs.

```

net_G = build_res_unet(n_input=1, n_output=2, size=256)
net_G.load_state_dict(torch.load("models/net_G_resnet18_model-19.pt",
↪ map_location=device))
model = MainModel(net_G=net_G)

```

```

train_model(model, train_dl, 20)
torch.save(model.state_dict(), "models/colorization2-epoch20.pt")

net_G = build_res_unet(n_input=1, n_output=2, size=224)
model = MainModel(net_G = net_G)
model.load_state_dict(torch.load("models/colorization2-epoch20.pt",
                                map_location=device))

```

model initialized with norm initialization

```

<All keys matched successfully>

my_paths = glob.glob("images/*.JPG")
my_dl = make_dataloaders(paths=my_paths, split='val',
                          n_workers=0)
my_iter = iter(my_dl)

data = next(my_iter)
fake_imgs, real_imgs = visualize(model, data, False)

```

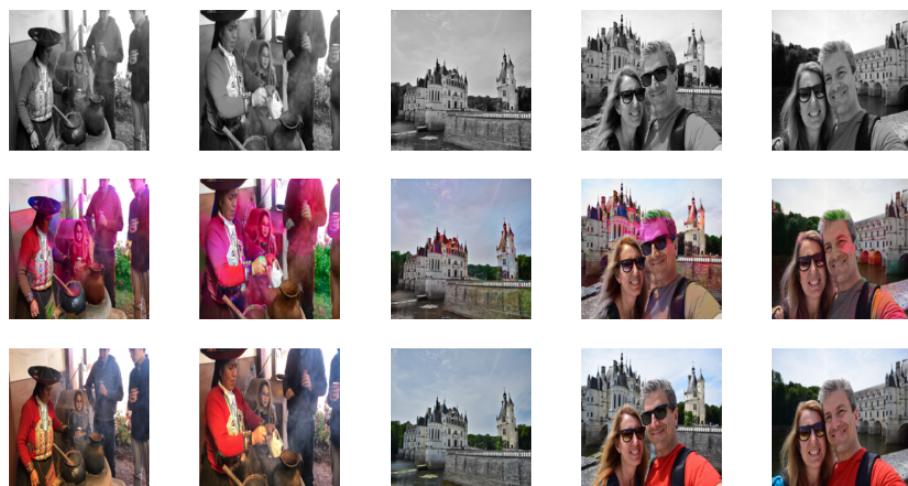


Figure 2: png

```

df = get_metrics(model, my_dl)

```

Table 1: ResNet18 metrics on validation dataset

Metric	a-channel	b-channel
0 R-square	0.9762806708182279	0.7944829539176959
1 Explained variance	0.9763600120897222	0.8047133044855702
2 Mean absolute error	0.022080948550583485	0.08126157218579984
3 Median absolute error	0.011479740269680015	0.053269025524179094
4 Mean squared error	0.0016633580971493562	0.014459893084060526
5 Sample size	2000	2000

```

df.to_latex(caption="ResNet18 metrics on personal
↪ photos", "Metrics of ResNet18 generator on dataset of personal
↪ photos"))

df

df = get_metrics(model, val_dl)

df.to_latex(caption="ResNet18 metrics on validation
↪ dataset", "Metrics of ResNet18 generator on validation
↪ dataset"))

df.to_pickle('results/ResNet18onValidation.pkl')

```

The results of the model with the ResNet18 generator are acceptable. However, many times they do not look natural because of an excessive use of colors by the generator that resulted in colorful blotches in the pictures.

In agreement with the visual inspection, the resulting metrics on the validation dataset showed that the network did a pretty good job at predicting the **a**-channel with over 97% of the variance of the channel predicted by the model with a very low mean squared error. However, the prediction on the **b**-channel was not as good with the model predicting only 80% of its variance.

Second approach: ViT as discriminator

```

#net_G = build_res_unet(n_input=1, n_output=2, size=256)
#net_G.load_state_dict(torch.load("models/net_G_resnet18_model-19.pt",
↪ map_location=device))
#net_D = build_visiontransformer()
#model = MainModel(net_G=net_G, net_D=net_D)
#model.load_state_dict(torch.load('models/model2-1.pt'))
#train_model(model, train_dl, 20, first_epoch=1)
#torch.save(net_G.state_dict(),
↪ "models/colorization3-epoch20.pt")

```

```

net_G = build_VTi_generator()
opt = optim.Adam(net_G.parameters(), lr=1e-4)
criterion = nn.L1Loss()
#pretrain_generator(net_G, train_dl, opt, criterion, 20)
#torch.save(net_G.state_dict(),
#            "models/net_G_ViT-20-pretraining.pt")
net_G.load_state_dict(torch.load("models/net_G_ViT-20-pretraining.pt",
                                 map_location=device))
model = MainModel(net_G = net_G, use_ViT_gen=True)
train_model(model, train_dl, 20)
torch.save(net_G.state_dict(),
            "models/colorization6-ViT-epoch20.pt")

```

```

net_G = build_VTi_generator()
#model = MainModel(net_G = net_G, use_ViT_gen=True)
#net_G = build_res_unet(n_input=1, n_output=2, size=224)
model = MainModel(net_G = net_G, use_ViT_gen=True)
# net_G.load_state_dict(torch.load("res18-unet.pt",
#                                 map_location=device))
#model = MainModel(net_G=net_G, net_D =
#                  build_visiontransformer())
net_G.load_state_dict(torch.load("models/colorization6-ViT-epoch20.pt",
                                 map_location=device))
test_paths = glob.glob("images/*.JPG")

my_dl = make_dataloaders(paths=test_paths, split='val',
                           n_workers=0) ## n_workers=0 in Windows

```

Some weights of the model checkpoint at google/vit-base-patch16-224-in21k were not used when initializing ViTForMaskedImageModeling from it.

- This IS expected if you are initializing ViTForMaskedImageModeling from the checkpoint of a different model.
- This IS NOT expected if you are initializing ViTForMaskedImageModeling from the checkpoint of the same model.

Some weights of ViTForMaskedImageModeling were not initialized from the model checkpoint at google/vit-base-patch16-224-in21k. You should probably TRAIN this model on a down-stream task to be able to use it for predictions.

model initialized with norm initialization

```

c:\users\fnoya\appdata\local\programs\python\python37\lib\site-packages\ipykernel_launcher.py
# Remove the CWD from sys.path while we load stuff.
df = get_metrics(model, val_dl)
df

```

0%| 0/125 [00:00<?, ?it/s]


```
c:\users\fnoya\appdata\local\programs\python\python37\lib\site-packages\skimage\_shared\util
    return func(*args, **kwargs)

Metric
a-channel
b-channel
0
R-square
0.9762806708182279
0.7944829539176959
1
Explained variance
0.9763600120897222
0.8047133044855702
2
Mean absolute error
0.022080948550583485
0.08126157218579984
3
Median absolute error
0.011479740269680015
0.053269025524179094
4
Mean squared error
0.0016633580971493562
0.014459893084060526
5
```

```
Sample size
```

```
2001
```

```
2001
```

```
test_iter = iter(my_dl)  
  
data = next(test_iter)  
fake_imgs, real_imgs = visualize(model, data, False)
```



Figure 3: png

```
def write_image(image:np.ndarray, image_path: str):  
    '''  
    Writes image from image path  
    Args:  
        image: RGB image of shape H x W x C, with float32 data  
        image_path: path to image  
  
    Returns:  
        RGB image of shape H x W x 3 in floating point format  
    '''  
    # read image and convert to RGB  
    bgr_image = (image[:, :, [2, 1, 0]]).astype(np.uint8)  
    cv2.imwrite(image_path, bgr_image)  
  
def rescale_img_with_colors (grayscale, pred):  
    h,w = grayscale.shape
```

```

f = cv2.resize(pred, (w,h))
f = cv2.cvtColor(f, cv2.COLOR_RGB2LAB)
grayscale = np.abs(grayscale*0.4).astype('uint8')
new = np.stack((grayscale, f[:,1], f[:,2]), axis=-1)
new = cv2.cvtColor(new, cv2.COLOR_LAB2RGB) * 255.
return new.astype('uint8')

test_paths = glob.glob("portraits/*.JPG")
my_dl = make_dataloaders(paths=test_paths, split='val',
                           n_workers=0)
my_iter = iter(my_dl)
i = 0
for data in tqdm(my_iter):
    fake_imgs, real_imgs = model_eval(model, data)
    for fimg in fake_imgs:
        imagefn = test_paths[i]
        i += 1
        original = cv2.imread(imagefn, cv2.IMREAD_GRAYSCALE)
        color = rescale_img_with_colors(original, fimg)
        write_image(color, "results/" + imagefn)
        fig, axes = plt.subplots(1, 3, figsize=(15,15))
        axes[0].imshow(original, cmap='gray')
        axes[1].imshow(fimg)
        axes[2].imshow(color)
        plt.show()

c:\users\fnoya\appdata\local\programs\python\python37\lib\site-packages\ipykernel_launcher.py
# Remove the CWD from sys.path while we load stuff.

```

0% | 0/1 [00:00<?, ?it/s]

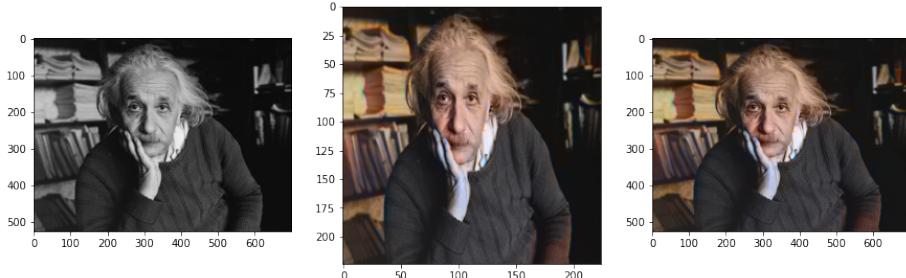


Figure 4: png

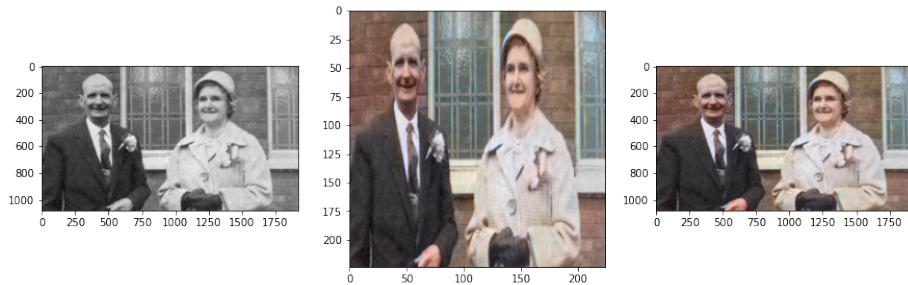


Figure 5: png

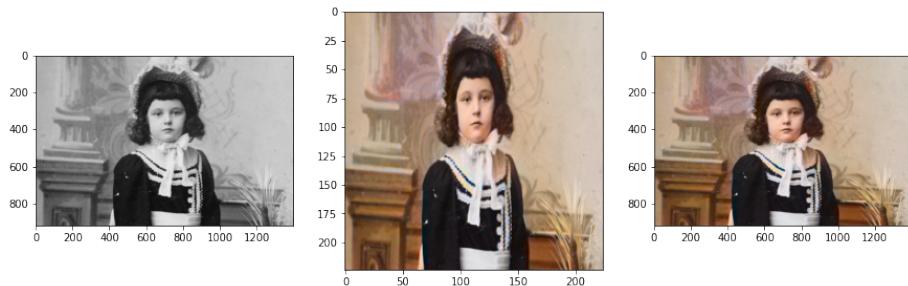


Figure 6: png

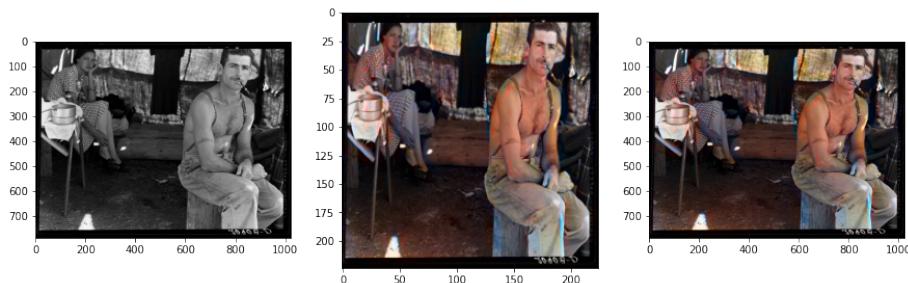


Figure 7: png

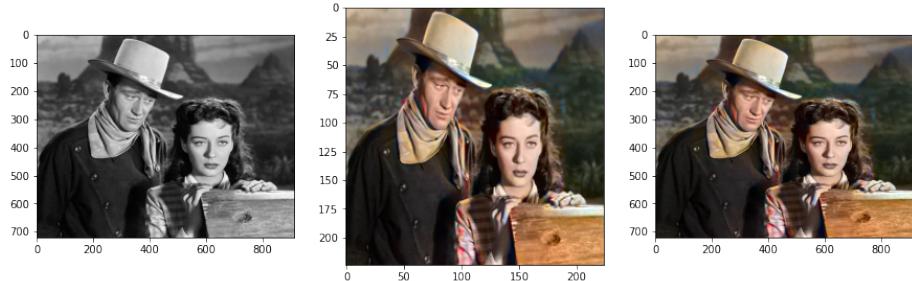


Figure 8: png

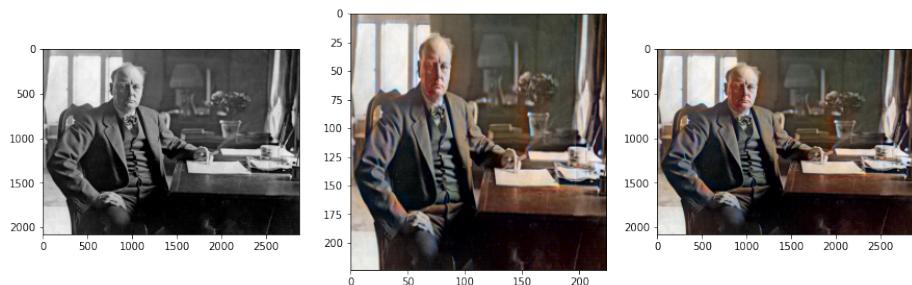


Figure 9: png

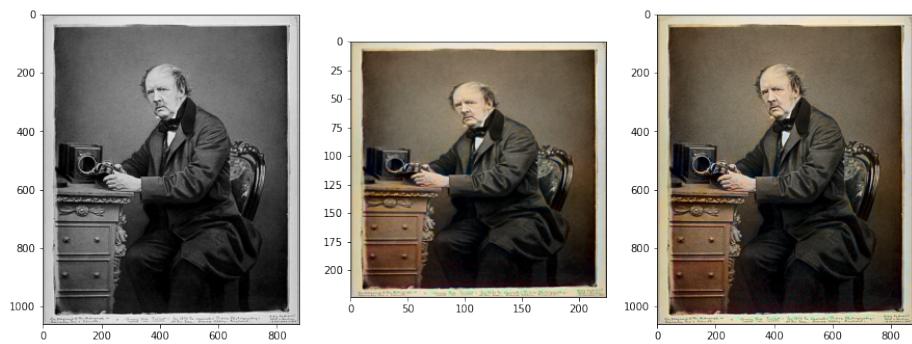


Figure 10: png

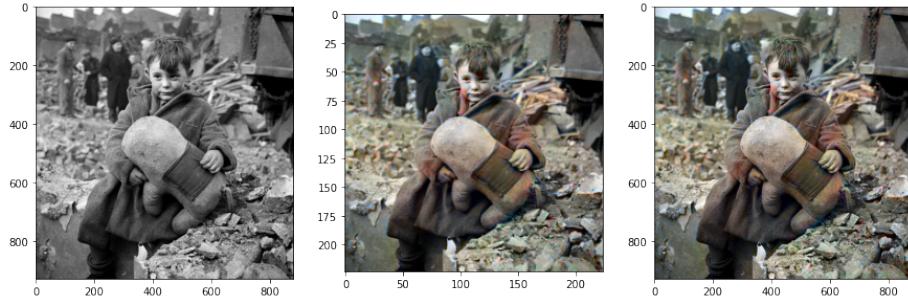


Figure 11: png

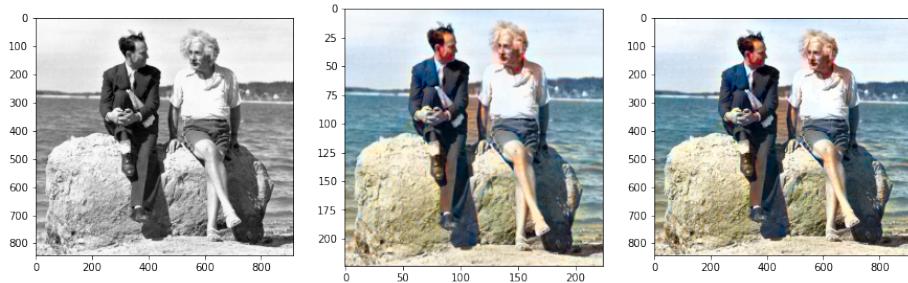


Figure 12: png



Figure 13: png



Figure 14: png

```
%matplotlib notebook
#%matplotlib widget

mask_coords = utils.specify_mask(img)
%matplotlib inline
```

If it doesn't get you to the drawing mode, then rerun this function again.

```
<IPython.core.display.Javascript object>

xs = mask_coords[0]
ys = mask_coords[1]
%matplotlib inline
import matplotlib.pyplot as plt
plt.figure()
fg_mask = utils.get_mask(ys, xs, img)
```

ValueError	Traceback (most recent call last)
	<pre><ipython-input-21-bb4a2f9fbf03> in <module> 4 import matplotlib.pyplot as plt 5 plt.figure() ----> 6 fg_mask = utils.get_mask(ys, xs, img)</pre>

```

~/SageMaker/CS445-Colorization-Project/utils.py in get_mask(ys, xs, img)
 115
 116 def get_mask(ys, xs, img):
--> 117     mask = poly2mask(ys, xs, img.shape[:2]).astype(int)
 118     fig = plt.figure()
 119     plt.imshow(mask, cmap='gray')

~/SageMaker/CS445-Colorization-Project/utils.py in poly2mask(vertex_row_coords, vertex_col_c
 4
 5 def poly2mask(vertex_row_coords, vertex_col_coords, shape):
----> 6     fill_row_coords, fill_col_coords = draw.polygon(vertex_row_coords, vertex_col_c
 7     mask = np.zeros(shape, dtype=np.bool)
 8     mask[fill_row_coords, fill_col_coords] = True

~/anaconda3/envs/mxnet_latest_p37/lib/python3.7/site-packages/skimage/draw/draw.py in polyg
497
498     """
--> 499     return _polygon(r, c, shape)
500
501

skimage/draw/_draw.pyx in skimage.draw._draw._polygon()

~/anaconda3/envs/mxnet_latest_p37/lib/python3.7/site-packages/numpy/core/_methods.py in _am
 32 def _amin(a, axis=None, out=None, keepdims=False,
 33             initial=_NoValue, where=True):
---> 34     return umr_minimum(a, axis, None, out, keepdims, initial, where)
 35
 36 def _sum(a, axis=None, dtype=None, out=None, keepdims=False,

ValueError: zero-size array to reduction operation minimum which has no identity

<Figure size 640x480 with 0 Axes>
fgModel = np.zeros((1, 65), dtype="float")
bgModel = np.zeros((1, 65), dtype="float")
fg_mask[fg_mask==1] = cv2.GC_PR_FGD
img = np.uint8(img*255)

```

```

fg_mask = np.uint8(fg_mask)
(mask, bgModel, fgModel) = cv2.grabCut(img, fg_mask, None,
    ↳ bgModel, \
                                fgModel, iterCount=5,
    ↳ mode=cv2.GC_INIT_WITH_MASK)

mask = np.where((mask==2) | (mask==0), 0, 1).astype('uint8')
fg_img = img*mask[:, :, np.newaxis]
plt.imshow(fg_img), plt.colorbar(), plt.show()

img_gray = cv2.cvtColor(cv2.cvtColor(img, cv2.COLOR_RGB2GRAY),
    ↳ cv2.COLOR_GRAY2RGB)
mask = mask[:, :, np.newaxis]
final_img = img_gray * (1 - mask) + fg_img * mask
plt.imshow(final_img), plt.show()

```

References

- <https://pyimagesearch.com/2020/07/27/opencv-grabcutforeground-segmentation-and-extraction/>
- https://docs.opencv.org/3.4/d8/d83/tutorial_py_grabcut.html
- <https://towardsdatascience.com/colorizing-black-white-images-with-u-net-and-conditional-gan-a-tutorial-81b2df111cd8>
- <https://doi.org/10.48550/arXiv.2106.06321>
- <https://machinelearningmastery.com/generative-adversarial-network-loss-functions/>
- [https://arxiv.org/abs/1711.10337 Are GANs Created Equal? A Large-Scale Study, 2018.](https://arxiv.org/abs/1711.10337)
- [https://arxiv.org/abs/1406.2661 Generative Adversarial Networks](https://arxiv.org/abs/1406.2661)
- [https://arxiv.org/pdf/1704.00028.pdf Improved Training of Wasserstein GANs](https://arxiv.org/pdf/1704.00028.pdf)
- [https://arxiv.org/abs/1611.04076 Least Squares Generative Adversarial Networks](https://arxiv.org/abs/1611.04076)

Colorizing black & white images with u-net and conditional GAN — a tutorial.

(n.d.). <https://towardsdatascience.com/colorizing-black-white-images-with-u-net-and-conditional-gan-a-tutorial-81b2df111cd8>.

Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., Uszkoreit, J., & Houlsby, N. (2020). An image is worth 16x16 words: Transformers for image recognition at scale. *CoRR, abs/2010.11929*. <https://arxiv.org/abs/2010.11929>

Fast.ai. (n.d.). <https://www.fast.ai/>.

Goodfellow, I. J., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., & Bengio, Y. (2014). *Generative adversarial networks*. <https://doi.org/10.48550/ARXIV.1406.2661>

Isola, P., Zhu, J.-Y., Zhou, T., & Efros, A. A. (2016). Image-to-image translation

- with conditional adversarial networks. *CoRR*, *abs/1611.07004*. <http://arxiv.org/abs/1611.07004>
- Lin, T.-Y., Maire, M., Belongie, S. J., Bourdev, L. D., Girshick, R. B., Hays, J., Perona, P., Ramanan, D., Dollár, P., & Zitnick, C. L. (2014). Microsoft COCO: Common objects in context. *CoRR*, *abs/1405.0312*. <http://arxiv.org/abs/1405.0312>
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., & Polosukhin, I. (2017). Attention is all you need. *CoRR*, *abs/1706.03762*. <http://arxiv.org/abs/1706.03762>
- Zhang, R., Isola, P., & Efros, A. A. (2016). Colorful image colorization. *CoRR*, *abs/1603.08511*. <http://arxiv.org/abs/1603.08511>